## CNN architecture (five convolutions, two poolings, global average pooling, linear classifier)

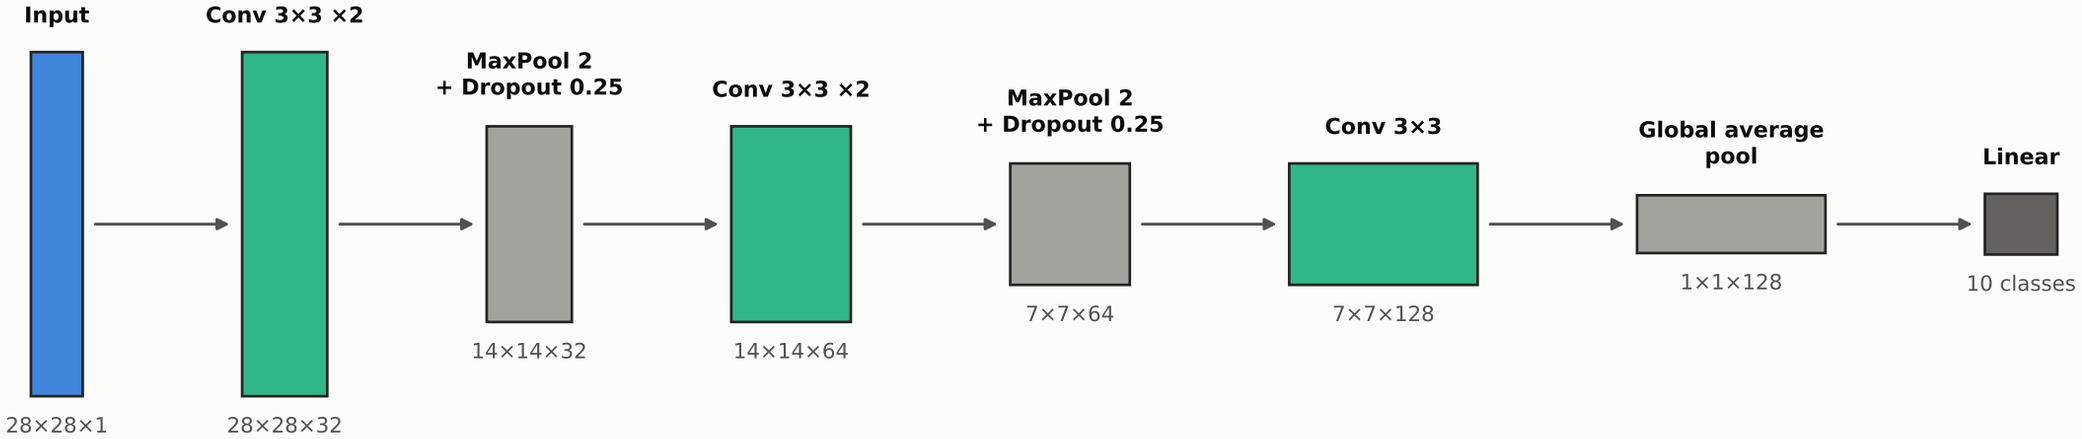


Every convolution: 3×3 kernel, padding 1, no bias, then BatchNorm and ReLU. Channels 32, 32, 64, 64, 128.

Fashion-MNIST input 28×28×1: 140,458 parameters. ciFAIR-10 input 32×32×3: 141,034 parameters (feature maps 32×32 → 16×16 → 8×8). Trained with Adam, one-cycle schedule, batch 128, 10 epochs, no augmentation.